

[RKE1](#) and [RKE2](#) clusters can be provisioned in Rancher using the built-in Harvester Node Driver. Harvester provides [load balancer](#) and Harvester cluster [storage passthrough](#) support to the guest Kubernetes cluster.

In this page we will learn:

- How to deploy the Harvester cloud provider in both RKE1 and RKE2 cluster.
- How to use the [Harvester load balancer](#).

## Backward Compatibility Notice

:::note

Please note a known backward compatibility issue if you're using the Harvester cloud provider version **v0.2.2** or higher. If your Harvester version is below **v1.2.0** and you intend to use newer RKE2 versions (i.e., `>= v1.26.6+rke2r1`, `v1.25.11+rke2r1`, `v1.24.15+rke2r1`), it is essential to upgrade your Harvester cluster to v1.2.0 or a higher version before proceeding with the upgrade of the guest Kubernetes cluster or Harvester cloud provider.

For a detailed support matrix, please refer to the **Harvester CCM & CSI Driver with RKE2 Releases** section of the official [website](#).

:::

## Deploying

### Prerequisites

- The Kubernetes cluster is built on top of Harvester virtual machines.
- The Harvester virtual machines run as guest Kubernetes nodes are in the same namespace.

:::info important

Each Harvester VM must have the `macvlan` kernel module, which is required for the `LoadBalancer` services of the **DHCP** IPAM mode.

To check if the kernel module is available, access the VM and run the following commands:

```
lsmod | grep macvlan
sudo modprobe macvlan
```

The kernel module is likely to be missing if the following occur:

- `$ lsmod | grep macvlan` does not produce output.
- `$ sudo modprobe macvlan` displays an error message similar to `modprobe: FATAL: Module macvlan not found in directory /lib/modules/5.14.21-150400.22-default`.

By default, the `macvlan` kernel module is not included in SUSE Linux Enterprise 15 Service Pack 4/5/6 minimal cloud images (see [Issue #6418](#)). Those images contain the `kernel-default-base` package, which includes only the base modules. However, the `macvlan` kernel driver becomes available when you install the `kernel-default` package.

To eliminate the need for manual intervention after the guest cluster is provisioned, build your own cloud images using the openSUSE Build Service (OBS). You must remove the `kernel-default-base` package and add the `kernel-default` package in the `Minimal.kiwi` file to ensure that the resulting cloud image includes the `macvlan` kernel module. For more information, see [Custom SUSE VM Images](#).

:::

## Deploying to the RKE1 Cluster with Harvester Node Driver

:::caution

Rancher Kubernetes Engine (RKE) will reach the end of its life on **July 31, 2025**. Harvester **v1.6.0** and later versions will not support RKE. Switching to RKE2, which provides a more secure and efficient environment, is recommended.

In-place upgrades are not an option, so you must [create new RKE2 clusters](#) and migrate the workloads from your existing RKE clusters (known as replatforming). For more information, see [RKE End of Life](#).

:::

When spinning up an RKE cluster using the Harvester node driver, you can perform two steps to deploy the `Harvester` cloud provider:

1. Select `Harvester(Out-of-tree)` option.
2. Install `Harvester Cloud Provider` from the Rancher marketplace.

## Deploying to the RKE2 Cluster with Harvester Node Driver

When spinning up an RKE2 cluster using the Harvester node driver, select the `Harvester` cloud provider. The node driver will then help deploy both the CSI driver and CCM automatically.

Starting with Rancher v2.9.0, you can configure a specific folder for cloud config data using the **Data directory configuration path** field.

## Manually Deploying to the RKE2 Cluster

1. Generate cloud config data using the script `generate_addon.sh`, and then place the data on every custom node (directory: `/etc/kubernetes/cloud-config`).

```
curl -sfL https://raw.githubusercontent.com/harvester/cloud-provider-harvester/master/deploy/generate_addon.sh | bash -s <serviceaccount name> <namespace>
```

:::note

The script depends on ``kubectl`` and ``jq`` when operating the Harvester cluster, and functions only when given access to the ``Harvester Cluster`` kubeconfig file.

You can find the ``kubeconfig`` file in one of the Harvester management nodes in the

`/etc/rancher/rke2/rke2.yaml` path. The server IP must be replaced with the VIP address.

Example of content:

```
```yaml
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: <redacted>
    server: https://127.0.0.1:6443
    name: default
# ...
```
```

You must specify the namespace in which the guest cluster will be created.

...

Example of output:

```
```yaml
##### cloud config #####
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: <CACERT>
    server: https://HARVESTER-ENDPOINT/k8s/clusters/local
    name: local
contexts:
- context:
    cluster: local
    namespace: default
    user: harvester-cloud-provider-default-local
    name: harvester-cloud-provider-default-local
current-context: harvester-cloud-provider-default-local
kind: Config
preferences: {}
users:
- name: harvester-cloud-provider-default-local
  user:
    token: <TOKEN>

##### cloud-init user data #####
write_files:
- encoding: b64
  content: <CONTENT>
  owner: root:root
  path: /etc/kubernetes/cloud-config
  permissions: '0644'
```
```

1. On the RKE2 cluster creation page, go to the **Cluster Configuration** screen and set the value of **Cloud Provider** to **External**.

2. Copy and paste the `cloud-init user data` content to **Machine Pools > Show Advanced > User Data**.

3. Add the `HelmChart` CRD for `harvester-cloud-provider` to **Cluster Configuration > Add-On Config > Additional Manifest**.

You must replace `<cluster-name>` with the name of your cluster.

```
apiVersion: helm.cattle.io/v1
kind: HelmChart
metadata:
  name: harvester-cloud-provider
  namespace: kube-system
spec:
  targetNamespace: kube-system
  bootstrap: true
  repo: https://raw.githubusercontent.com/rancher/charts/dev-v2.9
  chart: harvester-cloud-provider
  version: 104.0.2+up0.2.6
  helmVersion: v3
  valuesContent: |-
    global:
      cattle:
        clusterName: <cluster-name>
```

4. To create the load balancer, add the annotation `cloudprovider.harvesterhci.io/ipam:<dhcp|pool>` .

## Deploying to the RKE2 custom cluster (experimental)

1. Generate cloud config data using the script `generate_addon.sh` , and then place the data on every custom node (directory: `/etc/kubernetes/cloud-config` ).

```
curl -sL https://raw.githubusercontent.com/harvester/cloud-provider-harvester/master/deploy/generate_addon.sh | bash -s <serviceaccount name> <namespace>
```

:::note

The script depends on `kubectl` and `jq` when operating the Harvester cluster, and functions only when given access to the `Harvester Cluster` kubeconfig file.

You can find the `kubeconfig` file in one of the Harvester management nodes in the `/etc/rancher/rke2/rke2.yaml` path. The server IP must be replaced with the VIP address.

Example of content:

```
```yaml
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: <redacted>
    server: https://127.0.0.1:6443
    name: default
# ...
```
```

You must specify the namespace in which the guest cluster will be created.

:::

Example of output:

```
```yaml
##### cloud config #####
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: <CACERT>
    server: https://HARVESTER-ENDPOINT/k8s/clusters/local
    name: local
contexts:
- context:
    cluster: local
    namespace: default
    user: harvester-cloud-provider-default-local
    name: harvester-cloud-provider-default-local
current-context: harvester-cloud-provider-default-local
kind: Config
preferences: {}
users:
- name: harvester-cloud-provider-default-local
  user:
    token: <TOKEN>

##### cloud-init user data #####
write_files:
- encoding: b64
  content: <CONTENT>
  owner: root:root
```

```
path: /etc/kubernetes/cloud-config
permissions: '0644'
...
```

1. Create a VM in the Harvester cluster with the following settings:

- **Basics** tab: The minimum requirements are 2 CPUs and 4 GiB of RAM. The required disk space depends on the VM image.
- **Networks** tab: Specify a network name with the format `nic-<number>` .
- **Advanced Options** tab: Copy and paste the content of the **Cloud Config User Data** screen.

2. On the **Basics** tab of the **Cluster Configuration** screen, select **Harvester** as the **Cloud Provider** and then select **Create** to spin up the cluster.

1. On the **Registration** tab, perform the steps required to run the RKE2 registration command on the VM.

## Deploying to the K3s cluster with Harvester node driver (experimental)

When spinning up a K3s cluster using the Harvester node driver, you can perform the following steps to deploy the harvester cloud provider:

1. Use `generate_addon.sh` to generate cloud config.

```
curl -sL https://raw.githubusercontent.com/harvester/cloud-provider-
harvester/master/deploy/generate_addon.sh | bash -s <serviceaccount name>
<namespace>
```

The output will look as follows:

```
##### cloud config #####
apiVersion: v1
clusters:
- cluster:
  certificate-authority-data: <CACERT>
  server: https://HARVESTER-ENDPOINT/k8s/clusters/local
  name: local
contexts:
- context:
```

```

    cluster: local
    namespace: default
    user: harvester-cloud-provider-default-local
    name: harvester-cloud-provider-default-local
current-context: harvester-cloud-provider-default-local
kind: Config
preferences: {}
users:
- name: harvester-cloud-provider-default-local
  user:
    token: <TOKEN>

##### cloud-init user data #####
write_files:
- encoding: b64
  content: <CONTENT>
  owner: root:root
  path: /etc/kubernetes/cloud-config
  permissions: '0644'

```

2. Copy and paste the `cloud-init user data` content to **Machine Pools > Show Advanced > User Data**.
3. Add the following `HelmChart` `yaml` of `harvester-cloud-provider` to **Cluster Configuration > Add-On Config > Additional Manifest**.

```

apiVersion: helm.cattle.io/v1
kind: HelmChart
metadata:
  name: harvester-cloud-provider
  namespace: kube-system
spec:
  targetNamespace: kube-system
  bootstrap: true
  repo: https://charts.harvesterhci.io/
  chart: harvester-cloud-provider
  version: 0.2.2
  helmVersion: v3

```

4. Disable the `in-tree` cloud provider in the following ways:
  - o Click the `Edit as YAML` button.
  - o Disable `servicelb` and set `disable-cloud-controller: true` to disable the default K3s cloud controller.

```
machineGlobalConfig:
  disable:
    - servicelb
  disable-cloud-controller: true
```

- o Add `cloud-provider=external` to use the Harvester cloud provider.

```
machineSelectorConfig:
  - config:
      kubelet-arg:
        - cloud-provider=external
      protect-kernel-defaults: false
```

With these settings in place a K3s cluster should provision successfully while using the external cloud provider.

## Upgrade Cloud Provider

### Upgrade RKE2

The cloud provider can be upgraded by upgrading the RKE2 version. You can upgrade the RKE2 cluster via the Rancher UI as follows:

1. Click **≡ > Cluster Management**.
2. Find the guest cluster that you want to upgrade and select **> Edit Config**.
3. Select **Kubernetes Version**.
4. Click **Save**.

### Upgrade RKE/K3s

RKE/K3s upgrade cloud provider via the Rancher UI, as follows:

1. Click **≡ > RKE/K3s Cluster > Apps > Installed Apps**.
2. Find the cloud provider chart and select **> Edit/Upgrade**.
3. Select **Version**.
4. Click **Next > Update**.

:::info

The upgrade process for a [single-node guest cluster](#) may stall when the new `harvester-cloud-provider` pod is stuck in the *Pending* state. This issue is caused by a section in the `harvester-cloud-provider` deployment that describes the rolling update strategy. Specifically, the default value conflicts with the `podAntiAffinity` configuration in single-node clusters.

For more information, see [this GitHub issue comment](#). To address the issue, manually delete the old `harvester-cloud-provider` pod. You might need to do this multiple times until the new pod can be successfully scheduled.

:::



## Load Balancer Support

Once you've deployed the Harvester cloud provider, you can leverage the Kubernetes `LoadBalancer` service to expose a microservice within the guest cluster to the external world. Creating a Kubernetes `LoadBalancer` service assigns a dedicated Harvester load balancer to the service, and you can make adjustments through the `Add-on Config` within the Rancher UI.

### IPAM

Harvester's built-in load balancer offers both **DHCP** and **Pool** modes, and you can configure it by adding the annotation `cloudprovider.harvesterhci.io/ipam: $mode` to its corresponding service. Starting from Harvester cloud provider  $\geq$  v0.2.0, it also introduces a unique **Share IP** mode. A service shares its load balancer IP with other services in this mode.

- **DCHP:** A DHCP server is required. The Harvester load balancer will request an IP address from the DHCP server.
- **Pool:** An [IP pool](#) must be configured first. The Harvester load balancer controller will allocate an IP for the load balancer service following [the IP pool selection policy](#).
- **Share IP:** When creating a new load balancer service, you can re-utilize an existing load balancer service IP. The new service is referred to as a secondary service, while the currently chosen service is the primary one. To specify the primary service in the secondary service, you can add the annotation `cloudprovider.harvesterhci.io/primary-service: $primary-service-name`. However, there are two known limitations:
  - Services that share the same IP address can't use the same port.
  - Secondary services cannot share their IP with additional services.

:::note

Modifying the `IPAM` mode isn't allowed. You must create a new service if you intend to change the `IPAM` mode.

:::

## Health checks

Beginning with Harvester cloud provider v0.2.0, additional health checks of the `LoadBalancer` service within the guest Kubernetes cluster are no longer necessary. Instead, you can configure [liveness](#) and [readiness](#) probes for your workloads. Consequently, any unavailable pods will be automatically removed from the load balancer endpoints to achieve the same desired outcome.