

Fail to Deploy a Multi-node Cluster Due to Incorrect HTTP Proxy Setting

ISO Installation Without a Harvester Configuration File

Configure HTTP Proxy During Harvester Installation

In some environments, you configure [http-proxy](#) of [OS Environment](#) during Harvester installation.

Configure HTTP Proxy After First Node is Ready

After the first node is installed successfully, you login into the `Harvester GUI` to configure [http-proxy](#) of [Harvester System Settings](#).

Then you continue to add more nodes to the cluster.

One Node Becomes Unavailable

The issue you may encounter:

```
The first node is installed successfully.  
  
The second node is installed successfully.  
  
The third node is installed successfully.  
  
Then the second node changes to Unavailable state and cannot recover automatically.
```

Solution

When the nodes in the cluster do not use the HTTP Proxy to communicate with each other, after the first node is installed successfully, you need to configure [http-proxy.noProxy](#) against the CIDR used by those nodes.

For example, your cluster assigns IPs from CIDR `172.26.50.128/27` to nodes via DHCP/static setting, please add this CIDR to `noProxy`.

After setting this, you can continue to add new nodes to the cluster.

For more details, please refer to [Harvester issue 3091](#).

ISO Installation With a Harvester Configuration File

When a Harvester configuration file is used in ISO installation, please configure proper `http-proxy` in [Harvester System Settings](#).

PXE Boot Installation

When [PXE Boot Installation](#) is adopted, please configure proper `http-proxy` in [OS Environment](#) and [Harvester System Settings](#).

Generate a Support Bundle

Users can generate a support bundle in the Harvester GUI with the following steps:

- Click the `Support` link at the bottom-left of Harvester Web UI.

- Click **Generate Support Bundle** button.
- Enter a useful description for the support bundle and click **Create** to generate and download a support bundle.

:::note

Whenever you encounter an issue that may be related to workloads deployed in custom namespaces, configure the [support-bundle-namespaces](#) setting to include those namespaces as data sources. The support bundle only collects data from the configured namespaces.

For timeout errors, you can adjust the value of the [support-bundle-timeout](#) setting and then restart the support bundle generation process.

If you intend to use a non-default container image, you can configure the [support-bundle-image](#) setting.

:::

:::note

For information about collecting guest cluster logs and configuration files, see [Guest Cluster Log Collection](#).

:::

Manually Download and Retain a Support Bundle File

By default, a support bundle file is automatically generated, downloaded, and deleted after you click **Create** on the Harvester UI. However, you may want to retain a file for various reasons, including the following:

- You are unable to download the file because of network connectivity errors and other issues.
- You must use a previously generated file to troubleshoot issues (because generating a support bundle file takes time).
- You want to view information that only exists in a previously generated file.

Even if the file remains in the cluster, the Harvester UI does not provide a download link. Use the following workaround to generate, manually download, and retain a support bundle file:

Generate the File and Prevent Automatic Downloading

1. On the Harvester UI, click **Generate Support Bundle**.
2. When the progress indicator reaches 20% to 80%, close the browser tab to prevent automatic downloading of the generated file.
3. Retrieve a list of all support bundles in all namespaces using `kubectl`.

Example:

```
$ kubectl get supportbundle -A
NAMESPACE      NAME           ISSUE_URL      DESCRIPTION    AGE
harvester-system bundle-htl5f    sp1            3h43m
```

4. Retrieve the details of all existing support bundles using the command `kubectl get supportbundle -A -o yaml`.

Example:

```
$ kubectl get supportbundle -A -oyaml
apiVersion: v1
items:
- apiVersion: harvesterhci.io/v1beta1
  kind: SupportBundle
  metadata:
    creationTimestamp: "2024-02-02T11:18:09Z"
    generation: 5
    name: bundle-htl5f // resource name
    namespace: harvester-system
    resourceVersion: "1218311"
    uid: a3776373-05fe-4584-8a9a-baac3fa91bbf
  spec:
    description: sp1
    issueURL: ""
  status:
    conditions:
    - lastUpdateTime: "2024-02-02T11:18:38Z"
      status: "True"
      type: Initialized
    filename: supportbundle_db25ccb6-b52a-4f9d-97dd-db2df2b004d4_2024-02-02T11-18-10Z.zip // support bundle file name
    filesize: 8868712
    progress: 100 // 100 means successfully generated
    state: ready
```

The file is ready for downloading when the value of `progress` is "100" and the value of `state` is "ready".

Download the File

1. Create a download URL that includes the following information:

- [VIP](#) or DNS name
- Resource name of the file
- Parameter `?retain=true` : If you do not include this parameter, resources related to the support bundle are automatically deleted after the file is successfully downloaded.

Example:

```
https://{vip/dns-name}/v1/harvester/supportbundles/bundle-htl5f/download?retain=true
```

2. Download the file using either a command-line tool (for example, curl and wget) or a web browser.

Example:

```
curl -k https://{vip/dns-name}/v1/harvester/supportbundles/bundle-htl5f/download?retain=true -o sb2.zip
```

3. Verify that resources related to the support bundle were not deleted.

Example:

```
$ kubectl get supportbundle -A
```

NAMESPACE	NAME	ISSUE_URL	DESCRIPTION	AGE
harvester-system	bundle-htl5f		sp1	3h43m

(Optional) Delete the Related Resources

Retained support bundle files consume memory and storage resources. Each file is backed by a `supportbundle-manager-bundle*` pod in the `harvester-system` namespace, and the generated ZIP file is stored in the `/tmp` folder of the pod's memory-based filesystem.

Example:

```
$ kubectl get pods -n harvester-system
```

NAME	READY	STATUS	RESTARTS
supportbundle-manager-bundle-dtl2k-69dcc69b59-w64vl	1/1	Running	0

8m18s

You can delete the related resources using the following methods:

- Manual: Run the command `kubectl delete supportbundle -n {namespace} {resource-name}`. Deleting a support bundle object automatically deletes the pod that backs it.

Example:

```
$ kubectl delete supportbundle -n harvester-system bundle-htl5f
supportbundle.harvesterhci.io "bundle-htl5f" deleted

$ kubectl get supportbundle -A
No resources found
```

- Automatic: Harvester deletes the related resources based on how the following settings are configured:
 - [support-bundle-expiration](#): Defines the time allowed for retaining a support bundle file
 - [support-bundle-timeout](#): Defines the time allowed for generating a support bundle file

Manually Copy the Support Bundle File

You can run the command `kubectl cp` to copy the generated file from the backing pod.

Example:

```
kubectl cp harvester-system/supportbundle-manager-bundle-dtl2k-69dcc69b59-w64vl:/tmp/support-bundle-kit/supportbundle_db25ccb6-b52a-4f9d-97dd-db2df2b004d4_2024-02-02T11-18-10Z.zip bundle.zip
```

Manually Collect Data for Support Bundle

Harvester is unable to collect data and generate a support bundle when the node is inaccessible or not ready. The workaround is to run a script and compress the generated files.

1. Prepare the environment.

```
mkdir -p /tmp/support-bundle # ensure /tmp/support-bundle exists
echo 'JOURNALCTL="/usr/bin/journalctl -o short-precise"' > /tmp/common
export SUPPORT_BUNDLE_NODE_NAME=$(hostname)
```

2. Run the following commands:

- o Download the script: `curl -o collector-harvester https://raw.githubusercontent.com/rancher/support-bundle-kit/refs/heads/master/hack/collector-harvester`
- o Add executable permissions: `chmod +x collector-harvester`
- o Run the script: `./collector-harvester / /tmp/support-bundle`

3. Compress the files in `/tmp/support-bundle`, and then attach the archive to the related issue.

Known Limitations

- Replacing the backing pod prevents the support bundle file from being downloaded.

The support bundle file is stored in the `/tmp` folder of the pod's memory-based filesystem so it is removed when the pod is replaced during cluster and node rebooting, Kubernetes pod rescheduling, and other processes. After starting, the new pod regenerates the file but assigns a name that is different from the file name in the support bundle object.

Example:

1. A support bundle file is generated and retained.

```
$ kubectl get supportbundle -A -oyaml
apiVersion: v1
items:
- apiVersion: harvesterhci.io/v1beta1
  kind: SupportBundle
  metadata:
    creationTimestamp: "2024-02-06T11:01:19Z"
    generation: 5
    name: bundle-yr2vq
    namespace: harvester-system
    resourceVersion: "1583252"
    uid: eb8538cf-886b-4791-a7b0-dbc34dcee524
  spec:
    description: sp2
    issueURL: ""
  status:
    conditions:
    - lastUpdateTime: "2024-02-06T11:01:47Z"
      status: "True"
      type: Initialized
    filename: supportbundle_db25ccb6-b52a-4f9d-97dd-db2df2b004d4_2024-
02-06T11-01-20Z.zip // file is ready to download
    filesize: 7832010
    progress: 100
    state: ready
  kind: List
```

```
metadata:
  resourceVersion: ""
```

2. The backing pod restarts.

```
$ kubectl get pods -n harvester-system supportbundle-manager-bundle-yr2vq-c5484fbdf-9pz8d -oyaml
apiVersion: v1
kind: Pod
metadata:
  ...
  labels:
    app: support-bundle-manager
    pod-template-hash: c5484fbdf
    rancher/supportbundle: bundle-yr2vq
    name: supportbundle-manager-bundle-yr2vq-c5484fbdf-9pz8d
    namespace: harvester-system

  containerStatuses:
    - containerID:
containerid://ea82b63875c18a2b5b36afea6a47a99a5efd26464f94d401cde1727d175ef740
    ...
    name: manager
    ready: true
    restartCount: 1
    started: true
    state:
      running:
        startedAt: "2024-02-06T11:05:33Z" // pod's latest starting
timestamp, newer than the timestamp in support bundle's file name
```

3. The backing pod regenerates the file after it starts.

The name of the regenerated file is different from the file name recorded in the support bundle object.

```
$ kubectl exec -i -t -n harvester-system supportbundle-manager-bundle-yr2vq-c5484fbdf-9pz8d -- ls /tmp/support-bundle-kit -alh
total 2.2M
drwxr-xr-x 3 root root 4.0K Feb  6 11:05 .
-rw-r--r-- 1 root root 2.2M Feb  6 11:05 supportbundle_db25ccb6-b52a-4f9d-97dd-db2df2b004d4_2024-02-06T11-05-34Z.zip // different with above
file name
```

4. Attempts to download the regenerated file fail.

The following download URL cannot be used to access the regenerated file.

```
https://{vip/dns-name}/v1/harvester/supportbundles/bundle-yr2vq/download?retain=true .
```

- Retained support bundle files may affect system and node rebooting, node draining, and system upgrades.

Retained support bundle files are backed by pods in the `harvester-system` namespace. These pods are replaced during system and node rebooting, node draining, and system upgrades, consuming CPU and memory resources. Moreover, the regenerated files are very similar in content to the retained files, which means that storage resources are also unnecessarily consumed.

For more information, see [Issue 3383](#).

Access Embedded Rancher and Longhorn Dashboards

Available as of v1.1.0

You can now access the embedded Rancher and Longhorn dashboards directly on the `Support` page, but you must first go to the `Preferences` page and check the `Enable Extension developer features` box under `Advanced Features`.

:::note

We only support using the embedded Rancher and Longhorn dashboards for debugging and validation purposes. For Rancher's multi-cluster and multi-tenant integration, please refer to the docs [here](#).

:::

I can't access Harvester after I changed SSL/TLS enabled protocols and ciphers

If you changed [SSL/TLS enabled protocols and ciphers settings](#) and you no longer have access to Harvester GUI and API, it's highly possible that NGINX Ingress Controller has stopped working due to the misconfigured SSL/TLS protocols and ciphers. Follow these steps to reset the setting:

1. Following [FAQ](#) to SSH into Harvester node and switch to `root` user.

```
$ sudo -s
```

2. Editing setting `ssl-parameters` manually using `kubectl` :

```
# kubectl edit settings ssl-parameters
```

3. Deleting the line `value: ...` so that NGINX Ingress Controller will use the default protocols and ciphers.

```
apiVersion: harvesterhci.io/v1beta1
default: '{}'
kind: Setting
metadata:
  name: ssl-parameters
...
value: '{"protocols":"TLS99","ciphers":"WRONG_CIPHER"}' # <- Delete this line
```

4. Save the change and you should see the following response after exit from the editor:

```
setting.harvesterhci.io/ssl-parameters edited
```

You can further check the logs of Pod `rke2-ingress-nginx-controller` to see if NGINX Ingress Controller is working correctly.

Network interfaces are not showing up

You may need help finding the correct interface with a 10G uplink since the interface is not showing up. The uplink doesn't show up when the `ixgbe` module fails to load because an unsupported SFP+ module type is detected.

How to identify the issue with the unsupported SFP?

Execute the command `lspci | grep -i net` to see the number of NIC ports connected to the motherboard. By running the command `ip a`, you can gather information about the detected interfaces. If the number of detected interfaces is less than the number of identified NIC ports, then it's likely that the problem arises from using an unsupported SFP+ module.

Testing

You can perform a simple test to verify whether the unsupported SFP+ is the cause. Follow these steps on a running node:

1. Create the file `/etc/modprobe.d/ixgbe.conf` manually with the content:

```
options ixgbe allow_unsupported_sfp=1
```

1. Then run following command:

```
rmmod ixgbe && modprobe ixgbe
```

If the above steps are successful and the missing interface shows, we can confirm that the issue is an unsupported SFP+. However, the above test is not permanent and will be flushed out once rebooted.

Solution

Due to support issues, Intel restricts the types of SFPs used on their NICs. To make the above changes persistent, adding the following content to a [config.yaml](#) during installation is recommended.

```
os:
  write_files:
    - content: |
        options ixgbe allow_unsupported_sfp=1
      path: /etc/modprobe.d/ixgbe.conf
    - content: |
        name: "reload ixgbe module"
        stages:
          boot:
            - commands:
                - rmmod ixgbe && modprobe ixgbe
      path: /oem/99_ixgbe.yaml
```