

The following sections contain information useful in troubleshooting issues related to Harvester VM management.

VM Start Button is Not Visible

Issue Description

On rare occasions, the **Start** button is unavailable on the Harvester UI for VMs that are *Off*. Without that button, users are unable to start the VMs.

VM General Operations

On the Harvester UI, the **Stop** button is visible after a VM is created and started.

The **Start** button is visible after the VM is stopped.

When the VM is powered off from inside the VM, both the **Start** and **Restart** buttons are visible.

General VM Related Objects

A Running VM

The objects `vm` , `vmi` , and `pod` , which are all related to the VM, exist. The status of all three objects is `Running` .

```
# kubectl get vm
NAME    AGE      STATUS    READY
vm8     7m25s    Running   True

# kubectl get vmi
NAME    AGE    PHASE    IP            NODENAME    READY
vm8     78s    Running  10.52.0.199   harv41      True

# kubectl get pod
NAME                                READY    STATUS    RESTARTS    AGE
virt-launcher-vm8-tl46h            1/1     Running   0            80s
```

A VM Stopped Using the Harvester UI

Only the object `vm` exists and its status is `Stopped` . Both `vmi` and `pod` disappear.

```
# kubectl get vm
NAME    AGE    STATUS    READY
vm8     123m   Stopped   False
```

```
# kubectl get vmi
No resources found in default namespace.

# kubectl get pod
No resources found in default namespace.

#
```

A VM Stopped Using the VM's Poweroff Command

The objects `vm`, `vmi`, and `pod`, which are all related to the VM, exist. The status of `vm` is `Stopped`, while the status of `pod` is `Completed`.

```
# kubectl get vm
NAME    AGE    STATUS    READY
vm8     134m   Stopped   False

# kubectl get vmi
NAME    AGE    PHASE    IP            NODENAME    READY
vm8     2m49s  Succeeded  10.52.0.199   harv41      False

# kubectl get pod
NAME                                READY    STATUS    RESTARTS    AGE
virt-launcher-vm8-tl46h            0/1     Completed  0           2m54s
```

Issue Analysis

When the issue occurs, the objects `vm`, `vmi`, and `pod` exist. The status of the objects is similar to that of **A VM Stopped Using the VM's Poweroff Command**.

Example:

The VM `ocffm031v000` is not ready (`status: "False"`) because the virt-launcher pod is terminating (`reason: "PodTerminating"`).

```
- apiVersion: kubevirt.io/v1
  kind: VirtualMachine
  ...
  status:
    conditions:
      - lastProbeTime: "2023-07-20T08:37:37Z"
        lastTransitionTime: "2023-07-20T08:37:37Z"
        message: virt-launcher pod is terminating
        reason: PodTerminating
        status: "False"
        type: Ready
```

Similarly, the VMI (virtual machine instance) `ocffm031v000` is not ready (`status: "False"`) because the virt-launcher pod is terminating (`reason: "PodTerminating"`).

```
- apiVersion: kubevirt.io/v1
  kind: VirtualMachineInstance
  ...
  name: ocffm031v000
```

```

...
status:
  activePods:
    ec36a1eb-84a5-4421-b57b-2c14c1975018: aibfredg02
  conditions:
  - lastProbeTime: "2023-07-20T08:37:37Z"
    lastTransitionTime: "2023-07-20T08:37:37Z"
    message: virt-launcher pod is terminating
    reason: PodTerminating
    status: "False"
    type: Ready

```

On the other hand, the pod `virt-launcher-ocffm031v000-rrkss` is not ready (`status: "False"`) because the pod has run to completion (`reason: "PodCompleted"`).

The underlying container `0d7a0f64f91438cb78f026853e6bebf502df1bdeb64878d351fa5756edc98deb` is terminated, and the `exitCode` is 0.

```

- apiVersion: v1
  kind: Pod
  ...
  name: virt-launcher-ocffm031v000-rrkss
  ...
  ownerReferences:
  - apiVersion: kubevirt.io/v1
  ...
    kind: VirtualMachineInstance
    name: ocffm031v000
    uid: 8d2cf524-7e73-4713-86f7-89e7399f25db
    uid: ec36a1eb-84a5-4421-b57b-2c14c1975018
  ...
  status:
    conditions:
    - lastProbeTime: "2023-07-18T13:48:56Z"
      lastTransitionTime: "2023-07-18T13:48:56Z"
      message: the virtual machine is not paused
      reason: NotPaused
      status: "True"
      type: kubevirt.io/virtual-machine-unpaused
    - lastProbeTime: "null"
      lastTransitionTime: "2023-07-18T13:48:55Z"
      reason: PodCompleted
      status: "True"
      type: Initialized
    - lastProbeTime: "null"
      lastTransitionTime: "2023-07-20T08:38:56Z"
      reason: PodCompleted
      status: "False"
      type: Ready
    - lastProbeTime: "null"
      lastTransitionTime: "2023-07-20T08:38:56Z"
      reason: PodCompleted

```

```

    status: "False"
    type: ContainersReady
...
  containerStatuses:
  - containerID:
containerd://0d7a0f64f91438cb78f026853e6bebf502df1bdeb64878d351fa5756edc98deb
  image: registry.suse.com/suse/sles/15.4/virt-launcher:0.54.0-150400.3.3.2
  imageID: sha256:43bb08efdabb90913534b70ec7868a2126fc128887fb5c3c1b505ee6644453a2
  lastState: {}
  name: compute
  ready: false
  restartCount: 0
  started: false
  state:
    terminated:
      containerID:
containerd://0d7a0f64f91438cb78f026853e6bebf502df1bdeb64878d351fa5756edc98deb
      exitCode: 0
      finishedAt: "2023-07-20T08:38:55Z"
      reason: Completed
      startedAt: "2023-07-18T13:50:17Z"

```

A critical difference is that the `Stop` and `Start` actions appear in the `stateChangeRequests` property of `vm`.

```

status:
  conditions:
...
  printableStatus: Stopped
  stateChangeRequests:
  - action: Stop
    uid: 8d2cf524-7e73-4713-86f7-89e7399f25db
  - action: Start

```

Root Cause

The root cause of this issue is under investigation.

It is notable that the [source code](#) checks the status of `vm` and assumes that the object is starting. No `Start` and `Restart` operations are added to the object.

```

func (vf *vmformatter) canStart(vm *kubevirtv1.VirtualMachine, vmi
*kubevirtv1.VirtualMachineInstance) bool {
    if vf.isVMStarting(vm) {
        return false
    }
    ..
}

func (vf *vmformatter) canRestart(vm *kubevirtv1.VirtualMachine, vmi
*kubevirtv1.VirtualMachineInstance) bool {
    if vf.isVMStarting(vm) {

```

```

    return false
}
...
}

func (vf *vmformatter) isVMStarting(vm *kubevirtv1.VirtualMachine) bool {
    for _, req := range vm.Status.StateChangeRequests {
        if req.Action == kubevirtv1.StartRequest {
            return true
        }
    }
    return false
}

```

Workaround

To address the issue, you can force delete the pod using the command `kubectl delete pod virt-launcher-ocffm031v000-rrkss -n namespace --force`.

After the pod is successfully deleted, the `Start` button becomes visible again on the Harvester UI.

Related Issue

<https://github.com/harvester/harvester/issues/4659>

VM Stuck in Starting State with Error Message `not a device node`

Impacted versions: v1.3.0

Issue Description

Some VMs may fail to start and then become unresponsive after the cluster or some nodes are restarted. On the **Dashboard** screen of the Harvester UI, the status of the affected VMs is stuck at *Starting*.

Issue Analysis

The status of the pod related to the affected VM is `CreateContainerError`.

```

$ kubectl get pods
NAME                                READY   STATUS              RESTARTS   AGE
virt-launcher-vm1-w9bqs             0/2     CreateContainerError 0           9m39s

```

The phrase `failed to generate spec: not a device node` can be found in the following:

```

$ kubectl get pods -oyaml
apiVersion: v1
items:
  apiVersion: v1
  kind: Pod
  metadata:
  ...

```

```
containerStatuses:
- image: registry.suse.com/suse/sles/15.5/virt-launcher:1.1.0-150500.8.6.1
  imageID: ""
  lastState: {}
  name: compute
  ready: false
  restartCount: 0
  started: false
  state:
    waiting:
      message: 'failed to generate container
"50f0ec402f6e266870eafb06611850a5a03b2a0a86fdd6e562959719ccc003b5"
      spec: failed to generate spec: not a device node'
      reason: CreateContainerError
```

kubelet.log file:

```
file path: /var/lib/rancher/rke2/agent/logs/kubelet.log

E0205 20:44:31.683371    2837 pod_workers.go:1294] "Error syncing pod, skipping"
err="failed to \"StartContainer\" for \"compute\" with CreateContainerError: \"failed
t
o generate container
\\\"255d42ec2e01d45b4e2480d538ecc21865cf461dc7056bc159a80ee68c411349\\\" spec: failed
to generate spec: not a device node\"" pod="default/virt-laun
cher-caddytest-9tjzj" podUID=d512bf3e-f215-4128-960a-0658f7e63c7c
```

containerd.log file:

```
file path: /var/lib/rancher/rke2/agent/containerd/containerd.log

time="2024-02-21T11:24:00.140298800Z" level=error msg="CreateContainer within sandbox
\"850958f388e63f14a683380b3c52e57db35f21c059c0d93666f4fdaafe337e56\" for
&ContainerMetadata{Name:compute,Attempt:0,} failed" error="failed to generate
container \"5ddad240be2731d5ea5210565729cca20e20694e364e72ba14b58127e231bc79\" spec:
failed to generate spec: not a device node"
```

After adding debug information to `containerd`, it identifies the error message `not a device node` is upon the file `pvc-3c1b28fb-*`.

```
time="2024-02-22T15:15:08.557487376Z" level=error msg="CreateContainer within sandbox
\"d23af3219cb27228623cf8168ec27e64e836ed44f2b2f9cf784f0529a7f92e1e\" for
&ContainerMetadata{Name:compute,Attempt:0,} failed" error="failed to generate
container \"e4ed94fb5e9145e8716bcb87aae448300799f345197d52a617918d634d9ca3e1\" spec:
failed to generate spec: get device path:
/var/lib/kubelet/plugins/kubernetes.io/csi/volumeDevices/publish/pvc-3c1b28fb-683e-
4bf5-9869-c9107a0f1732/20291c6b-62c3-4456-be8a-fbeac118ec19 containerPath: /dev/disk-0
error: not a device node"
```

This is a CSI related file, but it is an empty file instead of the expected device file. Then the containerd denied the `CreateContainer` request.

```
$ ls /var/lib/kubelet/plugins/kubernetes.io/csi/volumeDevices/publish/pvc-3c1b28fb-683e-4bf5-9869-c9107a0f1732/ -alh
total 8.0K
drwxr-x--- 2 root root 4.0K Feb 22 15:10 .
-rw-r--r-- 1 root root 0 Feb 22 14:28 aa851da3-cee1-45be-a585-26ae766c16ca
-rw-r--r-- 1 root root 0 Feb 22 14:07 20291c6b-62c3-4456-be8a-fbeac118ec19
drwxr-x--- 4 root root 4.0K Feb 22 14:06 ..
-rw-r--r-- 1 root root 0 Feb 21 15:48 4333c9fd-c2c8-4da2-9b5a-1a310f80d9fd
-rw-r--r-- 1 root root 0 Feb 21 09:18 becc0687-b6f5-433e-bfb7-756b00deb61b

$file /var/lib/kubelet/plugins/kubernetes.io/csi/volumeDevices/publish/pvc-3c1b28fb-683e-4bf5-9869-c9107a0f1732/20291c6b-62c3-4456-be8a-fbeac118ec19
: empty
```

The output listed above directly contrasts with the following example, which shows the expected device file of a running VM.

```
$ ls /var/lib/kubelet/plugins/kubernetes.io/csi/volumeDevices/publish/pvc-732f8496-103b-4a08-83af-8325e1c314b7/ -alh
total 8.0K
drwxr-x--- 2 root root 4.0K Feb 21 10:53 .
drwxr-x--- 4 root root 4.0K Feb 21 10:53 ..
brw-rw---- 1 root root 8, 16 Feb 21 10:53 4883af80-c202-4529-a2c6-4e7f15fe5a9b
```

Root Cause

After the cluster or specific nodes are rebooted, the kubelet calls `NodePublishVolume` for the new pod without first calling `NodeStageVolume`. Moreover, the Longhorn CSI plugin bind mounts the regular file at the staging target path (previously used by the deleted pod) to the target path, and the operation is considered successful.

Workaround

Cluster level operation:

1. Find the backing pods of the affected VMs and the related Longhorn volumes.

```
$ kubectl get pods
NAME                                READY   STATUS              RESTARTS   AGE
virt-launcher-vm1-nxfm4             0/2     CreateContainerError 0           7m11s

$ kubectl get pvc -A
NAMESPACE      NAME                                STATUS    VOLUME
CAPACITY       ACCESS MODES  STORAGECLASS  AGE
default        vm1-disk-0-9gc6h                    Bound     pvc-f1798969-5b72-4d76-9f0e-64854af7b59c
1Gi           RWX          longhorn-image-fxsqr
7d22h
```

2. [Stop](#) the affected VMs from Harvester UI.

The VM may stuck in `Stopping`, continue the next step.

3. Delete the backing pods forcibly.

```
$ kubectl delete pod virt-launcher-vm1-nxfm4 --force
Warning: Immediate deletion does not wait for confirmation that the running
resource has been terminated. The resource may continue to run on the cluster
indefinitely.
pod "virt-launcher-vm1-nxfm4" force deleted
```

The VM is off now.

Node level operation, node by node:

1. [Cordon](#) a node.
2. Unmount all the affected Longhorn volumes in this node.

You need to ssh to this node and execute the `sudo -i umount path` command.

```
$ umount /var/lib/kubelet/plugins/kubernetes.io/csi/volumeDevices/pvc-f1798969-5b72-4d76-9f0e-64854af7b59c/dev/*
umount: /var/lib/kubelet/plugins/kubernetes.io/csi/volumeDevices/pvc-f1798969-5b72-4d76-9f0e-64854af7b59c/dev/4b2ab666-27bd-4e3c-a218-fb3d48a72e69: not
mounted.
umount: /var/lib/kubelet/plugins/kubernetes.io/csi/volumeDevices/pvc-f1798969-5b72-4d76-9f0e-64854af7b59c/dev/6aaf2bbe-f688-4dcd-855a-f9e2afa18862: not
mounted.
umount: /var/lib/kubelet/plugins/kubernetes.io/csi/volumeDevices/pvc-f1798969-5b72-4d76-9f0e-64854af7b59c/dev/91488f09-ff22-45f4-afc0-ca97f67555e7: not
mounted.
umount: /var/lib/kubelet/plugins/kubernetes.io/csi/volumeDevices/pvc-f1798969-5b72-4d76-9f0e-64854af7b59c/dev/bb4d0a15-737d-41c0-946c-85f4a56f072f: not
mounted.
umount: /var/lib/kubelet/plugins/kubernetes.io/csi/volumeDevices/pvc-f1798969-5b72-4d76-9f0e-64854af7b59c/dev/d2a54e32-4edc-4ad8-a748-f7ef7a2cacab: not
mounted.
```

3. [Uncordon](#) this node.
4. [Start](#) the affected VMs from harvester UI.

Wait some time, the VM will run successfully.

The newly generated csi file is an expected device file.

```
$ ls /var/lib/kubelet/plugins/kubernetes.io/csi/volumeDevices/publish/pvc-f1798969-5b72-4d76-9f0e-64854af7b59c/ -alh
...
brw-rw---- 1 root root 8, 64 Mar  6 11:47 7beb531d-a781-4775-ba5e-8773773d77f1
```

Related Issue

<https://github.com/harvester/harvester/issues/5109>

<https://github.com/longhorn/longhorn/issues/8009>

Virtual Machine IP Address Not Displayed

Issue Description

The **Virtual Machines** screen on the Harvester UI does not display the IP address of a newly created or imported virtual machine.

Issue Analysis

This issue usually occurs when the `qemu-guest-agent` package is not installed on the virtual machine. To determine if this is the root cause, check the status of the `VirtualMachineInstance` object.

```
$ kubectl get vmi -n <NAMESPACE> <NAME> -  
ojsonpath='{.status.interfaces[0].infoSource}'
```

The output does not contain the string `guest-agent` when the `qemu-guest-agent` package is not installed.

Workaround

You can [install the QEMU guest agent](#) by editing the virtual machine configuration.

1. On the Harvester UI, go to **Virtual Machines**.
2. Locate the affected virtual machine, and then select : > **Edit Config**.
3. On the **Advanced Options** tab, under **Cloud Config**, select **Install guest agent**.
4. Click **Save**.

However, cloud-init is run only once (when the virtual machine is started for the first time). To apply new **Cloud Config** settings, you must delete the cloud-init directory in the virtual machine.

```
$ sudo rm -rf /var/lib/cloud/*
```

After deleting the directory, you must restart the virtual machine so that cloud-init is run again and the `qemu-guest-agent` package is installed.

Related Issue

<https://github.com/harvester/harvester/issues/6644>