

General information

An **Upgrade** button appears on the **Dashboard** screen whenever a new Harvester version that you can upgrade to becomes available. For more information, see [Start an upgrade](#).

For air-gapped environments, see [Prepare an air-gapped upgrade](#).

Preventing Corruption of VM Images During Upgrade

:::caution

Before proceeding with the upgrade to Harvester **v1.4.0**, please make sure the **BackingImage** CRD is updated to the [Longhorn v1.7.2 version](#) beforehand.

If this step is skipped, it may lead to backing image corruption, as described in this [known Longhorn issue](#).

:::

To prevent the issue from occurring, you can manually update the `BackingImage` CRD before upgrading Harvester.

1. Patch the **Harvester managedchart** to avoid related errors and warnings.

```
kubectl patch managedchart harvester \
-n fleet-local \
--type='json' \
-p='[
  {
    "op": "add",
    "path": "/spec/diff/comparePatches/-",
    "value": {
      "apiVersion": "apiextensions.k8s.io/v1",
      "jsonPointers": ["/spec", "/metadata/annotations", "/metadata/labels",
"/status"],
      "kind": "CustomResourceDefinition",
      "name": "backingimages.longhorn.io"
    }
  }
]
```

1. Apply the **Longhorn v1.7.2** [BackingImage CRD](#).

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  annotations:
    controller-gen.kubebuilder.io/version: v0.15.0
  labels:
    app.kubernetes.io/name: longhorn
    app.kubernetes.io/instance: longhorn
    app.kubernetes.io/version: v1.7.2
    longhorn-manager: ""
  name: backingimages.longhorn.io
```

```

spec:
  conversion:
    strategy: Webhook
    webhook:
      clientConfig:
        service:
          name: longhorn-conversion-webhook
          namespace: longhorn-system
          path: /v1/webhook/conversion
          port: 9501
      conversionReviewVersions:
        - v1beta2
        - v1beta1
  group: longhorn.io
  names:
    kind: BackingImage
    listKind: BackingImageList
    plural: backingimages
    shortNames:
      - lhbi
    singular: backingimage
  scope: Namespaced
  versions:
    - additionalPrinterColumns:
        - description: The backing image name
          jsonPath: .spec.image
          name: Image
          type: string
        - jsonPath: .metadata.creationTimestamp
          name: Age
          type: date
      name: v1beta1
    schema:
      openAPIV3Schema:
        description: BackingImage is where Longhorn stores backing image object.
        properties:
          apiVersion:
            description: |-
              APIVersion defines the versioned schema of this representation of an
              object.
              Servers should convert recognized schemas to the latest internal value,
              and
              may reject unrecognized values.
              More info: https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#resources
            type: string
          kind:
            description: |-
              Kind is a string value representing the REST resource this object
              represents.
              Servers may infer this from the endpoint the client submits requests to.
              Cannot be updated.

```

```

    In CamelCase.
    More info: https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#types-kinds
    type: string
  metadata:
    type: object
  spec:
    x-kubernetes-preserve-unknown-fields: true
  status:
    x-kubernetes-preserve-unknown-fields: true
  type: object
served: true
storage: false
subresources:
  status: {}
- additionalPrinterColumns:
  - description: The system generated UUID
    jsonPath: .status.uuid
    name: UUID
    type: string
  - description: The source of the backing image file data
    jsonPath: .spec.sourceType
    name: SourceType
    type: string
  - description: The backing image file size in each disk
    jsonPath: .status.size
    name: Size
    type: string
  - description: The virtual size of the image (may be larger than file size)
    jsonPath: .status.virtualSize
    name: VirtualSize
    type: string
  - jsonPath: .metadata.creationTimestamp
    name: Age
    type: date
name: v1beta2
schema:
  openAPIV3Schema:
    description: BackingImage is where Longhorn stores backing image object.
    properties:
      apiVersion:
        description: |-
          APIVersion defines the versioned schema of this representation of an
object.
          Servers should convert recognized schemas to the latest internal value,
and
          may reject unrecognized values.
        More info: https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#resources
        type: string
      kind:
        description: |-

```

Kind is a string value representing the REST resource this object represents.

Servers may infer this from the endpoint the client submits requests to.

Cannot be updated.

In CamelCase.

More info: <https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#types-kinds>

type: string

metadata:

type: object

spec:

description: BackingImageSpec defines the desired state of the Longhorn backing image

properties:

checksum:

type: string

diskFileSpecMap:

additionalProperties:

properties:

evictionRequested:

type: boolean

type: object

type: object

diskSelector:

items:

type: string

type: array

disks:

additionalProperties:

type: string

description: Deprecated. We are now using DiskFileSpecMap to assign different spec to the file on different disks.

type: object

minNumberOfCopies:

type: integer

nodeSelector:

items:

type: string

type: array

secret:

type: string

secretNamespace:

type: string

sourceParameters:

additionalProperties:

type: string

type: object

sourceType:

enum:

- download

- upload

- export-from-volume

```

        - restore
        - clone
        type: string
    type: object
status:
  description: BackingImageStatus defines the observed state of the Longhorn
    backing image status
  properties:
    checksum:
      type: string
    diskFileStatusMap:
      additionalProperties:
        properties:
          lastStateTransitionTime:
            type: string
          message:
            type: string
          progress:
            type: integer
          state:
            type: string
        type: object
      nullable: true
      type: object
    diskLastRefAtMap:
      additionalProperties:
        type: string
      nullable: true
      type: object
    ownerID:
      type: string
    size:
      format: int64
      type: integer
    uuid:
      type: string
    virtualSize:
      description: Virtual size of image, which may be larger than physical
        size. Will be zero until known (e.g. while a backing image is
uploading)
      format: int64
      type: integer
    type: object
  type: object
served: true
storage: true
subresources:
  status: {}

```

3. Start the upgrade process.

Known issues

1. A VM with a container disk can't be migrated which makes the upgrade stuck in pre-drain status

:::tip

Manually stop the VMs to continue the upgrade process.

:::

When upgrading from v1.3.2 to v1.4.0, the upgrade process may become stuck if a VM with a container disk cannot be migrated. There is some limitation of live migration.

For more information, see [Issue #7005](#).

2. Upgrade stuck on waiting for Harvester bundle

When upgrading from v1.3.2 to v1.4.0, the upgrade process may become stuck on waiting for the Harvester bundle to become ready. This issue is caused by a race condition when the Fleet agent (`fleet-agent`) is redeployed.

The following error messages indicate that the issue exists.

```
> kubectl get bundles -n fleet-local
NAME                                BUNDLEDEPLOYMENTS-READY  STATUS
mcc-harvester                      0/1
ErrApplied(1) [Cluster fleet-local/local: encountered 2 deletion errors. First is:
admission webhook "validator.harvesterhci.io" denied the request: Internal error
occurred: no route match found for DELETE /v1, Kind=Secret harvester-
system/sh.helm.release.v1.harvester.v2]
mcc-harvester-crd                  0/1
ErrApplied(1) [Cluster fleet-local/local: admission webhook
"validator.harvesterhci.io" denied the request: Internal error occurred: no route
match found for DELETE /v1, Kind=Secret harvester-
system/sh.helm.release.v1.harvester-crd.v1]
```

You can run the following script to fix the issue.

```
#!/bin/bash

patch_fleet_bundle() {
    local bundleName=$1
    local generation=$(kubectl get -n fleet-local bundle ${bundleName} -o
jsonpath='{.spec.forceSyncGeneration}')
    local new_generation=$((generation+1))
    patch_manifest="$(mktemp)"
    cat > "$patch_manifest" <<EOF
{
  "spec": {
    "forceSyncGeneration": $new_generation
```

```

    }
}
EOF
    echo "patch bundle to new generation: $new_generation"
    kubectl patch -n fleet-local bundle ${bundleName} --type=merge --patch-file
$patch_manifest
    rm -f $patch_manifest
}

echo "removing harvester validating webhook"
kubectl delete validatingwebhookconfiguration harvester-validator

for bundle in mcc-harvester-crd mcc-harvester
do
    patch_fleet_bundle ${bundle}
done

echo "removing longhorn services"
kubectl delete svc longhorn-engine-manager -n longhorn-system --ignore-not-
found=true
kubectl delete svc longhorn-replica-manager -n longhorn-system --ignore-not-
found=true

```

3. Upgrade stuck on waiting for Fleet

When upgrading from v1.3.2 to v1.4.0, the upgrade process may become stuck on waiting for Fleet to become ready. This issue is caused by a race condition when Rancher is redeployed.

Check the Harvester logs and Fleet history for the following indicators:

- The manifest pod is stuck in the `deployed` status.
- The upgrade is pending with a chart version that has been deployed.

Example:

```

> kubectl logs -n harvester-system -l harvesterhci.io/upgradeComponent=manifest
wait helm release cattle-fleet-system fleet fleet-104.0.2+up0.10.2 0.10.2 deployed

> helm history -n cattle-fleet-system fleet
REVISION      UPDATED              STATUS      CHART
APP VERSION   DESCRIPTION
26            Tue Dec 10 03:09:13 2024    superseded  fleet-
103.1.5+up0.9.5 0.9.5                Upgrade complete
27            Sun Dec 15 09:26:54 2024    superseded  fleet-
103.1.5+up0.9.5 0.9.5                Upgrade complete
28            Sun Dec 15 09:27:03 2024    superseded  fleet-
103.1.5+up0.9.5 0.9.5                Upgrade complete
29            Mon Dec 16 05:57:03 2024    deployed    fleet-
103.1.5+up0.9.5 0.9.5                Upgrade complete
30            Mon Dec 16 05:57:13 2024    pending-upgrade fleet-
103.1.5+up0.9.5 0.9.5                Preparing upgrade

```

You can run the following command to fix the issue.

```
helm rollback fleet -n cattle-fleet-system <last-deployed-revision>
```

4. Upgrade will start over again unexpectedly after clicking the "Dismiss it" button

When you use Rancher to upgrade Harvester, the Rancher UI displays a dialog with a button labeled "Dismiss it". Clicking this button may result in the following issues:

- The `status` section of the `harvesterhci.io/v1beta1/upgrade` CR is cleared, causing the loss of all important information about the upgrade.
- The upgrade process starts over again unexpectedly.

This issue affects Rancher v2.10.x, which uses v1.0.2, v1.0.3, and v1.0.4 of the [Harvester UI Extension](#). All Harvester UI versions are not affected. The issue will be fixed in Harvester UI Extension v1.0.5 and v1.5.0.

To avoid this issue, perform either of the following actions:

- Use the Harvester UI to upgrade Harvester. Clicking the "Dismiss it" button on the Harvester UI does not result in unexpected behavior.
- Instead of clicking the button on the Rancher UI, run the following command against the cluster:

```
kubectl -n harvester-system label upgrades -l harvesterhci.io/latestUpgrade=true  
harvesterhci.io/read-message=true
```

Related issue:

- [\[BUG\] upgrade controller does not handle read-message well due to UI menu Dismiss it wipes upgrade CR's status](#)